Bachelor's in Computer Science and Engineering

Bachelor's in Computer Science and Engineering and Business Administration

Programming Assignment 2: File System

Operating Systems Design

David del Rio Astorga
Javier López Gómez

Version 1.3

Contents

1 Introduction

The aim of this programming assignment is to develop a simplified userspace file system using the C language (ISO C11) on the VM provided for guided labs and/or *guernika*, whose storage will be backed by a file in the host operating system.

In this assignment, the student will:

1. Design the architecture of the file system, its control structures (e.g. i-nodes, superblock or file descriptors) and algorithms in order to comply with the requirements and features defined in this document.

2. Implement the former design in C language, and develop client programs relying on the interface of this file system to access and modify the files within.

3. Justify the design and implementation internals of the file system, in tight relation to the theoretical concepts covered in the Operating Systems Design subject. This is particularly critical for implementation dependant or undefined behaviours.

4. Build a validation plan for the resulting design against the given requirements.

5. Summarise and discuss the former points in a written report.

All of these items will be taken into consideration for the evaluation of the assignment.

2 Evaluation

The final mark of this project is obtained as follows:

- **Design** (*4 points*).
- **Validation plan** (*2 point*).
- **Implementation** (*4 points*).

3 Initial Code

The file *dssoo_fs.zip* contains baseline code and utilities for your convenience. Inside you will find a folder containing the following:

- **Files that CANNOT be modified**

  1. `include/blocks_cache.h`: declaration of the functions to read or write blocks from/to the disk. This includes the `bread` and `bwrite` functions to read and write, respectively, a full block from the device given its block number. It is strictly **FORBIDDEN** to access the disk in a different manner.
  2. `blocks_cache.c`: implementation of the functions to read or write blocks from/to the disk.
  3. `include/filesystem.h`: declaration of the interface with the file system. This interface has to be implemented in `filesystem.c`.
  4. `create_disk.c`: creates `disk.dat` files with different sizes.
  5. `Makefile`: used by the build tool to compile the program. Use `make` to compile the program and `make clean` to remove compiled files.

- **Files to be completed by the student**

  1. `authors.txt`: file that should contain NIAs of group members, one per line.
  2. `filesystem.c`: implementation of the interface to use the file system. The student must use this file as starting point in order to develop the file system requested. Any additional function developed by the student must be implemented here, and it must be exclusively internal to this file.
  3. `include/metadata.h`: definition of the structures, data types and constants defined by the student in order to implement the file system.

     Additionally, you can make use of the `bitmap_getbit()` and `bitmap_setbit()` functions to handle bit arrays of arbitrary length, and could be useful to track used blocks.

     `bitmap_getbit(bitmap_, i_)`: gets the state of the i-th bit in the bitmap referenced to by `bitmap_` (of type `char *`).

     `bitmap_setbit(bitmap_, i_, val_)`: sets to `val_` the state of the i-th bit in the bitmap referenced to by `bitmap_` (of type `char *`).

     The following listing gives an idea of their usage:

     ```
     char bitmap[2];  // array of lenth = 16 bits
     int val = bitmap_getbit(bitmap, 7);
     printf("%d\n", val); // Value of bit 7 = 0
     bitmap_setbit(bitmap, 7, 1);
     val = bitmap_getbit(bitmap, 7);
     printf("%d\n", val); // Value of bit 7 = 1
     ```

  4. `include/auxiliary.h`: definition of auxiliary functions for the core functions defined in `filesystem.h`. These functions must be implemented in `filesystem.c` along with the core functionality, and they cannot be used outside the later (i.e. this file does not extend the interface of the file system).
  5. `test.c`: includes a minimal set of tests to check some features of the file system. This file should be extended with further tests created by the student to validate the features of the file system exhaustively, and check all possible errors detailed in the interface.

It is strictly **FORBIDDEN** to modify the signature of the functions in the files `blocks_cache.h`, `crc.h`, and `filesystem.h` (i.e. name, parameters, return type).

Further information on the return values of the provided functions can be found in the documentation in the header files. The student must consult this information before proceeding to the implementation of the design.

4 Specification of Functionality

The student has to design and implement from scratch a file system able to manage an emulated storage device (`disk.dat`).

The following sections define in detail the client interface for the whole requested functionality. These functions are declared in `filesystem.h`, and have to be implemented within `filesystem.c`.

4.1 Device Management

```
int mkFS(long deviceSize)
```

- **Behaviour:** generates the proper file system structure in a storage device, as designed by the student.
- **Parameters:**
  - `deviceSize` – Size of the disk to be formatted, in *bytes*.
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case of error. Trying to create a file system which exceeds the storage capacity limit of the device is considered an error.

```
int mountFS(void)
```

- **Behaviour:** this is the first FS-related operation to be executed by a client program in order to interact with the files. This function mounts the simulated device –*disk.dat*–, hence allocating and setting up all the structures and variables necessary to use the FS.
- **Parameters:** none.
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case of error.

```
int unmountFS(void)
```

- **Behaviour:** this is the last FS-related operation to be executed by a client program, as it unmounts the simulated device *disk.dat*. This function frees all the structures and variables used by the FS.
- **Parameters:** none.
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case of error.

4.2 File Management

```
int createFile(char *fileName)
```

- **Behaviour:** performs all the necessary changes in the file system in order to create a new, empty file.
- **Parameters:**
  - `fileName` – Name of the file to create.
    (e.g. `/Movies/Mission_Impossible/poster`).
- **Return:**
  - `0` – The execution is correct.
  - `-1` – The file cannot be created because it already exists in the file system.
  - `-2` – In case of other errors.

```
int removeFile(char *fileName)
```

- **Behaviour:** performs all the necessary changes in the file system in order to remove a file.
- **Parameters:**
  - `fileName` – Name of the file to remove.
    (e.g. `/Movies/Mission_Impossible/intro_song`).
- **Return:**
  - `0` – The execution is correct.
  - `-1` – The file cannot be removed because it does not exist in the file system.
  - `-2` – In case of other errors.

```
int openFile(char *fileName)
```

- **Behaviour:** opens an existing file in the file system and initialises its seek pointer to the beginning of the file.
- **Parameters:**
  - `fileName` – Name of the file to open.
    (e.g. `/Movies/Mission_Impossible/poster2`).
- **Return:**
  - The file descriptor of the opened file.
  - `-1` – The file cannot be opened because it does not exist in the file system.
  - `-2` – In case of other errors.

```
int closeFile(int fileDescriptor)
```

- **Behaviour:** closes an opened file.
- **Parameters:**
  - `fileDescriptor` – File descriptor of the file to close.
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case of error.

4.3 Interacting with Files

```
int readFile(int fileDescriptor, void *buffer, int numBytes)
```

- **Behaviour:** reads as many bytes as indicated from a given file descriptor, starting from the seek pointer of the file, and placing the data into the provided buffer. This function increments the seek pointer of the file as many bytes as read from the file.
- **Parameters:**
  - `fileDescriptor` – File descriptor of the file to read from.
  - `buffer` – Buffer that will store the read data after the execution of the function.
  - `numBytes` – Number of bytes to read from the file.
- **Return:**
  - Number of bytes correctly read. Consider that if the remaining amount of bytes to be read from the file is lower than the `numBytes` parameter, the function will return the former. In particular, reading when there are no remaining bytes will result in a return value of 0.
  - `-1` – In case of error.

```
int writeFile(int fileDescriptor, void *buffer, int numBytes)
```

- **Behaviour:** modifies the file given by a file descriptor, starting from the seek pointer of the file. It stores as many bytes as indicated from a user-provided buffer. This function increments the seek pointer of the file as many bytes as written into the file.
- **Parameters:**
  - `fileDescriptor` – File descriptor of the file to write into.
  - `buffer` – Data to be written.
  - `numBytes` – Number of bytes to write from the buffer.
- **Return:**

- Number of bytes correctly written. Trying to write past the maximum file size will not be considered an error, and the number of bytes written in the block will be returned. Hence, trying to write when the pointer is at EOF will return 0 bytes written.
- `-1` – In case of error.

```
int lseekFile(int fileDescriptor, int whence, long offset)
```

- **Behaviour:** modifies the value of the file seek pointer according to a given reference and offset.
- **Parameters:**
  - `fileDescriptor` – File descriptor.
  - `whence` – Constant value acting as reference for the seek operation. This could be:
    - `FS_SEEK_CUR` – Current position of the seek pointer.
    - `FS_SEEK_BEGIN` – Beginning of the file.
    - `FS_SEEK_END` – End of the file.
  - `offset` – Number of bytes to displace the seek pointer from the `FS_SEEK_CUR` position. This value can be either positive or negative, although it will never allow positioning the seek pointer outside the limits of the file. If `whence` is set to `FS_SEEK_BEGIN` or `FS_SEEK_END`, the file pointer must be set to this position, regardless of the offset value.
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case of error.

4.4 Directory management

```
int mkDir(char *path)
```

- **Behaviour:** performs all the necessary changes to the filesystem in order to create an empty directory.
- **Parameters:**
  - `path` – Full path of the directory that shall be created.
    (e.g. `/Movies/Mission_Impossible`).
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case the directory already exists in the filesystem.
  - `-2` – In case of other errors.

```
int rmDir(char *path)
```

- **Behaviour:** performs all the necessary changes to the filesystem in order to remove a directory and its contents.
- **Parameters:**
  - `path` – Full path of the directory that is to be removed. (e.g. `/Movies/Mission_Impossible`).
- **Return:**
  - `0` – The execution is correct.
  - `-1` – In case the directory does not exists in the filesystem.
  - `-2` – In case of other errors.

4.5 Interacting with directories

```
int lsDir(char *path, int inodesDir[10], char namesDir[10][33])
```

- **Behaviour:** loads directory contents into the provided arrays.
- **Parameters:**
  - `path` – Full path of the directory that shall be read. (e.g. `/Movies/Mission_Impossible`).
  - `inodesDir` – Array to be filled with the inodes used by the elements contained in the given directory.
  - `namesDir` – Array to be filled with the names of the elements contained in the given directory, following the same order as `inodesDir`.
- **Return:**
  - Number of elements in the directory.
  - `-1` – In case the directory does not exist in the filesystem.
  - `-2` – In case of other errors.

4.6 Undefined Behaviours

Any behaviour or architectural feature not specified in this document is subject to the student's judgement, thus considered a *design decision*. Since the objective of this assignment is to design a file system, design decisions have a key role in assessing the student's knowledge on the topic.

The student has to **clearly state in the report** those characteristics which are *implementation dependant* or *undefined*, i.e. out of the scope of this document, but not in conflict with this specification or the requirements defined in Sec. 5. Therefore, the student must describe these problems, ambiguities, or obscure issues, and justify the selected approach.

5 Specification of Requirements

The student must design and implement the file system to attain the functionality described in Sec. 4. In addition, the student shall comply with the requirements defined in this section. It is highly recommended that the student verifies the compliance with those.

5.1 Functional Requirements

F1 The file system shall support the following core functionality:

- **F1.1** Create a file system (function `mkFS`).
- **F1.2** Mount a file system (function `mountFS`).
- **F1.3** Unmount a file system (function `unmountFS`).
- **F1.4** Create a file within the file system (function `createFile`).
- **F1.5** Remove an existing file from the file system (function `removeFile`).
- **F1.6** Open an existing file (function `openFile`).
- **F1.7** Close an opened file (function `closeFile`).
- **F1.8** Read from an opened file (function `readFile`).
- **F1.9** Write to an opened file (function `writeFile`).
- **F1.10** Modify the position of the seek pointer (function `lseekFile`).
- **F1.11** Create a new directory in the filesystem (function `mkDir`).
- **F1.12** Remove an existing directory in the filesystem (function `rmDir`).
- **F1.13** List contents of an existing directory (function `lsDir`).

F2 Every time a file is opened, its seek pointer will be reset to the beginning of the file.

F3 Metadata shall be updated after any write operation in order to properly reflect any modification in the file system.

F4 The whole contents of a file could be read by means of several `read` operations.

F5 A file could be modified by means of `write` operations.

F6 The file system can be created on partitions of the device smaller than its maximum size.

F7 The character '`/`' will be used as path separator. Therefore, it cannot be part of the name of a directory entry.

F8 The root directory path is considered to be'`/`'. It may contain both files (e.g. '`/photo1`') and/or directories, that in turn might be nested (e.g. '`/music/soundtracks/`').

5.2 Non-functional Requirements

NF1 The maximum number of files in the file system will never be higher than 40.

NF2 The maximum number of elements (files and directories) contained in a direcrtory will never be higher than 10.

NF3 The maximum length of a file or directory name will be 32 characters.

NF4 The maximum length of a full path to an existing file (including its name) will be 132 characters, including path separators ('/'). In case a path names a directory, this limit will be 99 characters, including path separators ('/').

NF5 The maximum depth of a directory hierarchy will not be greater that 3.

NF6 The maximum size of a file will be 1 block.

NF7 The file system block size will be 2048 bytes.

NF8 Metadata shall persist between `unmount` and `mount` operations.

NF9 The file system will be used on disks from 50 KiB to 10 MiB.

NF10 The size of on-disk filesystem metadata shall be minimized.

NF11 The implementation shall not waste available resources.

5.3 Implementation Requirements

I1 The submitted code has to work on the *guernika* server and/or on the provided VMs for the guided labs. You must compile and test your code in either one before submitting[1].

I2 The code must compile with the flags `-Werror` and `-Wall` in order to be evaluated. Programs that do not compile will have a grade of zero.

I3 The files `blocks_cache.h`, `blocks_cache.c`, `filesystem.h`, `create_disk.c` and `Makefile` cannot be modified[2].

I4 The storage device will be emulated in a single file named `disk.dat`. Using other files is forbidden.

I5 The implementation of the file system shall comply with the client interface defined in Sec. 4.

I7 The student shall declare auxiliary functions in `auxiliary.h`.

I8 Both the client interface of the file system and the auxiliary functions defined by the student shall be implemented in `filesystem.c`.

I9 The student shall declare auxiliary data types, constants and structures in `metadata.h`.

5.4 Documentation Requirements

D1 Each function in the code shall be properly commented, emphasising obscure implementation details or student-defined behaviour. Programs without comments will have their grade penalised.

D2 A minimal report must include:

D2.1 A title page with the name of the authors and their student identification numbers.

D2.2 An index.

[1] Contact your lab teacher if you have issues or inquiries regarding this platform.

[2] Contact your lab teacher in case you find any bugs or issues you consider not working as intended **BEFORE** modifying the code yourself.

D2.3 The detailed design of the file system, including assumptions, data structures, algorithms and optimisations.

D2.4 A high-level description of the core functionality.

D2.5 The design of the test plan to verify the functionality of an implementation of the design.

D2.6 Conclusions, describing the main problems found throughout the assignment, and how they have been solved. Optionally, here you can include personal comments regarding the project.

D3 *Every* design decision must be properly justified and contextualised with respect to the theoretical concepts of the subject.

D4 Do not include any source code in the report.

D5 Do not include any source code in screenshot form in the report.

D6 Each case in the test plan must specify its objective, procedure, input, and expected output. Optionally, you can also include whether your implementation passes the case.

D7 Every page must be numbered, except the cover.

D8 Text must be justified.

D9 The report must not exceed 15 pages, including cover, figures, tables, and references. Any content exceeding this limit will be ignored.

5.5 Submission Requirements

S1 You must submit your work using Aula Global. The deadline will be indicated in the Aula Global assignment, and notified accordingly. Submitting the assignments by mail is not allowed without prior authorization.

S2 Submission must be done separately for the code and the report.

S3 The report must be submitted as a PDF file through Turnitin. Only PDF files will be graded.

S4 The submission must comply with the following naming conventions, where AAAAAAAAA, BBBBBBBBB, and CCCCCCCCC are the student identification numbers of the students in each group:

S4.1 Report: `dssoo_p2_AAAAAAAAA_BBBBBBBBB_CCCCCCCCC.pdf`

S4.2 Code: `dssoo_p2_AAAAAAAAA_BBBBBBBBB_CCCCCCCC.zip`

S5 The compressed file shall contain the following files:

S5.1 Implementation of the filesystem: `filesystem.c`

S5.2 Implementation of the test plan: `test.c`

S5.3 Full header file folder: `include`

S6 A maximum of three students is allowed per submission.

S7 All the requirements should be checked before submission.

6 Other Considerations

In addition, the following items need to be considered by the student:

- Provided two file systems, both passing all tests, the better design will achieve higher scores.
- The report is a significant part of the grade given the design-oriented nature of this assignment. Do not neglect its quality.
- Avoid duplicated tests that target the same code paths with equivalent input parameters. The test plan is graded according to its coverage, rather than the number of tests.
- A warning-free compilation does not guarantee that the implementation fulfills the functional requirements. Please test and debug your code properly to assess its behaviour.
- Students are expected to submit original work. In case plagiarism is detected between two assignments both groups will fail the continuous evaluation. Additional administrative charges of academic misconduct may be applied.
- You can use the following command to create the compressed file with the source code:

```
&> zip -r dssoo_p3_AAAAAAAAA_BBBBBBBBB_CCCCCCCCC.zip filesystem.c
test.c include/*
```